Multiple Issue Introduction

[Adapted from Mary Jane Irwin for
Computer Organization and Design,
Patterson & Hennessy, © 2005, UCB]

Other handouts
To handout next time

Review: Pipeline Hazards

❑ Structural hazards

  ● Design pipeline to eliminate structural hazards

❑ Data hazards – read before write

  ● Use data forwarding inside the pipeline

  ● For those cases that forwarding won't solve (e.g., load-use) include hazard hardware to insert stalls in the instruction stream

❑ Control hazards – `beq,bne,j,jr,jal`

  ● Stall – hurts performance

  ● Move decision point as early in the pipeline as possible – reduces number of stalls at the cost of additional hardware

  ● Delay decision (requires compiler support) – not feasible for deeper pipes requiring more than one delay slot to be filled

  ● Predict – with even more hardware, can reduce the impact of control hazard stalls even further if the branch prediction (BHT) is correct and if the branched-to instruction is cached (BTB)

Extracting Yet *More* Performance

❑ Two options:

- ● Increase the depth of the pipeline to increase the clock rate – superpipelining
- ● Fetch (and execute) more than one instructions at one time (expand every pipeline stage to accommodate multiple instructions) – multiple-issue

❑ Launching multiple instructions per stage allows the instruction execution rate, CPI, to be less than 1

- ● So instead we use IPC: instructions per clock cycle
  - - E.g., a 6 GHz, four-way multiple-issue processor can execute at a peak rate of 24 billion instructions per second with a best case CPI of 0.25 or a best case IPC of 4
  - - If the datapath has a five stage pipeline, how many instructions are active in the pipeline at any given time?

Answer to the last question is, obviously, twenty!

CPI clock cycles per instruction

Superpipelined Processors

❑ Increase the depth of the pipeline leading to shorter clock cycles (and more instructions "in flight" at one time)

  ● The higher the degree of superpipelining, the more forwarding/ hazard hardware needed, the more pipeline latch overhead (i.e., the pipeline latch accounts for a larger and larger percentage of the clock cycle time), and the bigger the clock skew issues (i.e., because of faster and faster clocks)

Superpipelined vs Superscalar

❑ Superpipelined processors have longer instruction latency than the SS processors which can degrade performance in the presence of true dependencies

❑ Superscalar processors are more susceptible to resource conflicts – but we can fix this with hardware !

Instruction latency: Anzahl Takte, die ein Befehl braucht.

Instruction vs Machine Parallelism

❑ Instruction-level parallelism (ILP) of a program – a measure of the average number of instructions in a program that a processor *might* be able to execute at the same time

  ● Mostly determined by the number of true (data) dependencies and procedural (control) dependencies in relation to the number of other instructions

❑ Data-level parallelism (DLP)

```
DO  I = 1  TO  100
    A[I] = A[I] + 1
CONTINUE
```

❑ Machine parallelism of a processor – a measure of the ability of the processor to take advantage of the ILP of the program

  ● Determined by the number of instructions that can be fetched and executed at the same time

❑ To achieve high performance, need *both* ILP and machine parallelism

The example shown has lots of data parallelism but almost no ILP. A good compiler could help with loop unrolling. If completely unrolled, and if had 100 arithmetic/addressing units and 100 memory ports, could achieve a speedup of 100 over a scalar processor

Multiple-Issue Processor Styles

❑ Static multiple-issue processors (aka VLIW)

- Decisions on which instructions to execute simultaneously are being made statically (at compile time by the compiler)
- E.g., Intel Itanium and Itanium 2 for the IA-64 ISA – EPIC (Explicit Parallel Instruction Computer)

❑ Dynamic multiple-issue processors (aka superscalar)

- Decisions on which instructions to execute simultaneously are being made dynamically (at run time by the hardware)
- E.g., IBM Power 2, Pentium 4, MIPS R10K, HP PA 8500

VLIW: Very long instruction word (weil ein Befehl mehrere Befehle enthält, die dann parallel ausgeführt werden)

Multiple-Issue Datapath Responsibilities

❑ Must handle, with a combination of hardware and software fixes, the fundamental limitations of

- ● Storage (data) dependencies – aka data hazards
    - Limitation more severe in a SS/VLIW processor due to (usually) low ILP
- ● Procedural dependencies – aka control hazards
    - Ditto, but even more severe
    - Use dynamic branch prediction to help resolve the ILP issue
- ● Resource conflicts – aka structural hazards
    - A SS/VLIW processor has a much larger number of potential resource conflicts
    - Functional units may have to arbitrate for result buses and register-file write ports
    - Resource conflicts can be eliminated by duplicating the resource or by pipelining the resource

Pipelining is much less expensive than duplicating

ILP = instruction level parallelism

SS = super scalar

VLIW = very long instruction word

Instruction Issue and Completion Policies

❑ Instruction-issue – initiate execution

  ● Instruction lookahead capability – fetch, decode and issue instructions beyond the current instruction

❑ Instruction-completion – complete execution

  ● Processor lookahead capability – complete issued instructions beyond the current instruction

❑ Instruction-commit – write back results to the RegFile or D$ (i.e., change the machine state)

In-order issue with in-order completion

In-order issue with out-of-order completion

Out-of-order issue with out-of-order completion

Completion == execution here

Out-of-order execution = an instruction blocked from executing does not prevent the following instructions from executing.

Instructions are executed in a different order from the one they were fetched

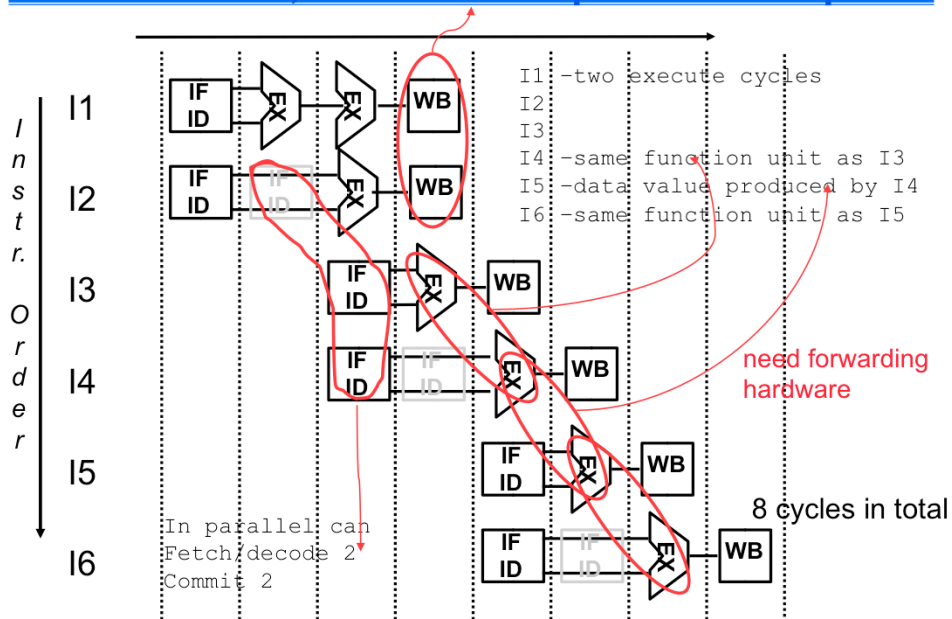Today all dynamically scheduled pipelines use in-order commit.

In-Order Issue with In-Order Completion

❑ Simplest policy is to issue instructions in exact program order and to complete them in the same order they were fetched (i.e., in program order)

❑ Example:

● Assume a pipelined processor that can fetch and decode two instructions per cycle, that has three functional units (a single cycle adder, a single cycle shifter, and a two cycle multiplier), and that can complete (and write back) two results per cycle

● And an instruction sequence with the following characteristics

```
I1 - needs two execute cycles (a multiply)
I2
I3
I4 - needs the same function unit as I3
I5 - needs data value produced by I4
I6 - needs the same function unit as I5
```
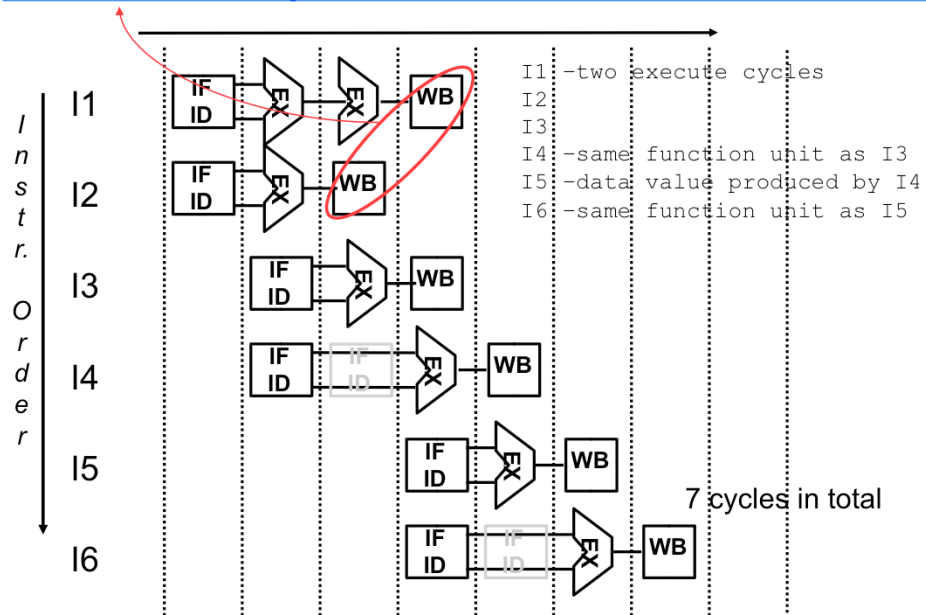
In-Order Issue, In-Order Completion Example

I1 -two execute cycles
I2
I3
I4 -same function unit as I3
I5 -data value produced by I4
I6 -same function unit as I5

need forwarding hardware

8 cycles in total

In parallel can
Fetch/decode 2
Commit 2

I3 can nicht früher starten, da nur 2 commit gleichzeitig möglich sind.

In-Order Issue with Out-of-Order Completion

❑ With out-of-order completion, a later instruction may complete before a previous instruction

- Out-of-order completion is used in single-issue pipelined processors to improve the performance of long-latency operations such as divide

❑ When using out-of-order completion instruction issue is stalled when there is a resource conflict (e.g., for a functional unit) or when the instructions ready to issue need a result that has not yet been computed

IOI-OOC Example


```
I1 -two execute cycles
I2
I3
I4 -same function unit as I3
I5 -data value produced by I4
I6 -same function unit as I5
```

7 cycles in total

For lecture

what is the best case time and what is the worst case time for this sequence?

Handling Output Dependencies

❑ There is one more situation that stalls instructions issuing with IOI-OOC, assume

I1 – writes to R3
I2 – writes to R3
I5 – reads R3

- ● If the I1 write occurs after the I2 write, then I5 reads an incorrect value for R3

- ● I2 has an output dependency on I1 – write before write

  - The issuing of I2 would have to be stalled if its result might later be overwritten by a previous instruction (i.e., I1) that takes longer to complete – the stall happens before instruction issue

❑ While IOI-OOC yields higher performance, it requires more dependency checking hardware

- ● Dependency checking needed to resolve both read before write and write before write
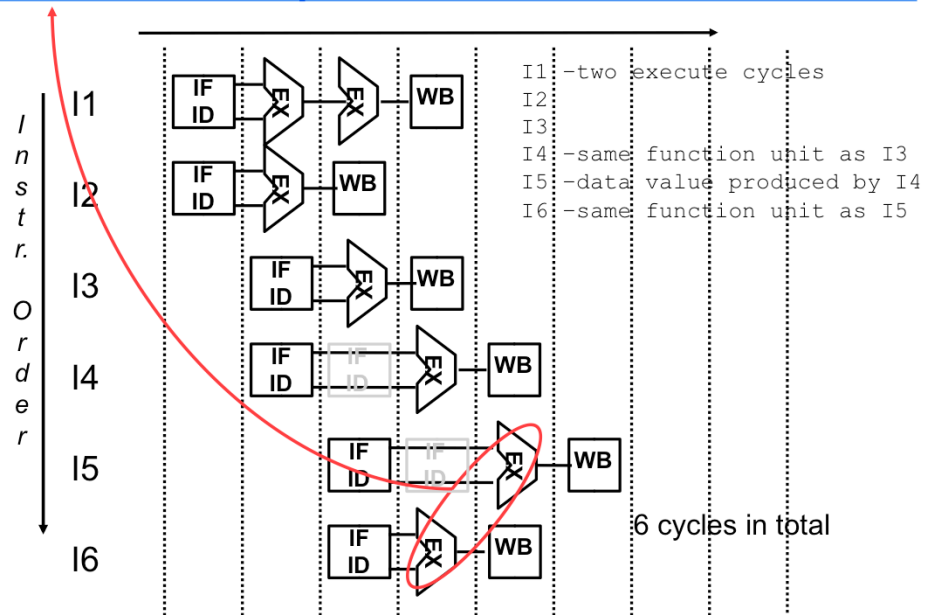
Need arbitration and more ports because there probably are not enough to satisfy all instructions that can complete simultaneously

Out-of-Order Issue with Out-of-Order Completion

❑ With in-order issue the processor stops decoding instructions whenever a decoded instruction has a resource conflict or a data dependency on an issued, but uncompleted instruction

- The processor is not able to *look beyond* the conflicted instruction even though more downstream instructions might have no conflicts and thus be issueable

❑ Fetch and decode instructions *beyond* the conflicted one, store them in an instruction buffer (as long as there's room), and flag those instructions in the buffer that don't have resource conflicts or data dependencies

❑ Flagged instructions are then issued from the buffer without regard to their program order

An instruction being in the instruction bufferonly implies that the processor has sufficient information about the instruction to know whether or not it can be issued

Instr. Order

I1
I2
I3
I4
I5
I6

IF ID — EX — EX — WB
IF ID — EX — WB
IF ID — EX — WB
IF ID — IF ID — EX — WB
IF ID — IF ID — EX — WB
IF ID — EX — WB

```
I1 -two execute cycles
I2
I3
I4 -same function unit as I3
I5 -data value produced by I4
I6 -same function unit as I5
```

6 cycles in total

Rechnerarchitektur   15

Assume that I4 is using a different functional unit than I6

Antidependencies

❏ With OOI *also* have to deal with data antidependencies – when a later instruction (that completes earlier) produces a data value that destroys a data value used as a source in an earlier instruction (that issues later)

```
 R3  := R3 * R5          True data dependency
 R4  := R3 + 1           Output dependency
 R3  := R5 + 1           Antidependency
```

❏ The constraint is similar to that of true data dependencies, except *reversed*

- Instead of the later instruction using a value (not yet) produced by an earlier instruction (read before write), the later instruction produces a value that destroys a value that the earlier instruction (has not yet) used (write before read)

Dependencies Review

❑ Each of the three data dependencies
  ● True data dependencies (read before write)
  ● Antidependencies (write before read)
  ● Output dependencies (write before write) } storage conflicts

  manifests itself through the use of registers (or other storage locations)

❑ True dependencies represent the flow of data and information through a program

❑ Anti- and output dependencies arise because the limited number of registers mean that programmers reuse registers for different computations

❑ When instructions are issued out-of-order, the correspondence between registers and values breaks down and the values *conflict* for registers

Storage conflicts, like other resource conflicts can be reduced or eliminated by duplicating the troublesome resource

Storage Conflicts and Register Renaming

❑ Storage conflicts can be reduced (or eliminated) by increasing or duplicating the troublesome resource

● Provide additional registers that are used to reestablish the correspondence between registers and values

- Allocated dynamically by the hardware in SS processors

❑ Register renaming – the processor renames the original register identifier in the instruction to a new register (one not in the visible register set)

```
R3  := R3 * R5              R3b := R3a * R5a
R4  := R3 + 1       ⟹      R4a := R3b + 1
R3  := R5 + 1              R3c := R5a + 1
```

● The hardware that does renaming assigns a "replacement" register from a pool of free registers and releases it back to the pool when its value is superseded and there are no outstanding references to it